

Team 1 Housing Data Analysis

Section 1 Data Importing and Preprocessing

Importing Libraries

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Loading data

We used a dataset of house_sales in a comma separated value format. We imported the .csv file using the pandas .read_csv method to load the data into a dataframe

```
In [ ]: #creating the dataframe by importing the csv file
housing_data = pd.read_csv('house_sales.csv')
#displaying the first 5 rows of the dataframe
print(housing_data.head())
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900.0	3.0	1.00	1180.0	
1	6414100192	20141209T000000	538000.0	3.0	2.25	2570.0	
2	5631500400	20150225T000000	180000.0	2.0	1.00	770.0	
3	2487200875	20141209T000000	604000.0	4.0	3.00	1960.0	
4	1954400510	20150218T000000	510000.0	3.0	2.00	1680.0	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	\
0	5650.0	1.0	0	0	...	7	1180	0	
1	7242.0	2.0	0	0	...	7	2170	400	
2	10000.0	1.0	0	0	...	6	770	0	
3	5000.0	1.0	0	0	...	7	1050	910	
4	8080.0	1.0	0	0	...	8	1680	0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	

	sqft_lot15
0	5650
1	7639
2	8062
3	5000
4	7503

[5 rows x 21 columns]

Now we are going to get the dimensions of the newly imported dataframe named housing_data

```
In [ ]: #dataframe dimensions
print(f"The dimensions of the initial dataframe are {housing_data.shape}")
```

The dimensions of the initial dataframe are (21613, 21)

Next we will get a sense for the columns in the dataframe and their data types

```
In [ ]: #printing columns in dataframe
print(housing_data.columns)
```

```
Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
       'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
       'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
       'lat', 'long', 'sqft_living15', 'sqft_lot15'],
      dtype='object')
```

```
In [ ]: #printing the data types of the columns
print(housing_data.dtypes)
```

```
id          int64
date        object
price       float64
bedrooms    float64
bathrooms   float64
sqft_living float64
sqft_lot    float64
floors      float64
waterfront  int64
view        int64
condition   int64
grade       int64
sqft_above  int64
sqft_basement int64
yr_built    int64
yr_renovated int64
zipcode     int64
lat         float64
long        float64
sqft_living15 int64
sqft_lot15  int64
dtype: object
```

Next we will view summary statistics of the initial dataset to get a feel for it.

```
In [ ]: #printing the descriptive statistics of the dataframe
print(housing_data.describe())
```

	id	price	bedrooms	bathrooms	sqft_living
count	2.161300e+04	2.161300e+04	20479.000000	20545.000000	20503.000000
mean	4.580302e+09	5.400881e+05	3.372821	2.113507	2081.073697
std	2.876566e+09	3.671272e+05	0.930711	0.768913	915.043176
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.500000	1430.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1920.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	12050.000000

	sqft_lot	floors	waterfront	view	condition
count	2.056900e+04	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.517982e+04	1.494309	0.007542	0.234303	3.409430
std	4.148617e+04	0.539989	0.086517	0.766318	0.650743
min	5.200000e+02	1.000000	0.000000	0.000000	1.000000
25%	5.040000e+03	1.000000	0.000000	0.000000	3.000000
50%	7.620000e+03	1.500000	0.000000	0.000000	3.000000
75%	1.070800e+04	2.000000	0.000000	0.000000	4.000000
max	1.651359e+06	3.500000	1.000000	4.000000	5.000000

	grade	sqft_above	sqft_basement	yr_built	yr_renovated
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	7.656873	1788.390691	291.509045	1971.005136	84.402258
std	1.175459	828.090978	442.575043	29.373411	401.679240
min	1.000000	290.000000	0.000000	1900.000000	0.000000
25%	7.000000	1190.000000	0.000000	1951.000000	0.000000
50%	7.000000	1560.000000	0.000000	1975.000000	0.000000
75%	8.000000	2210.000000	560.000000	1997.000000	0.000000
max	13.000000	9410.000000	4820.000000	2015.000000	2015.000000

	zipcode	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	98077.939805	47.560053	-122.213896	1986.552492	12768.455652
std	53.505026	0.138564	0.140828	685.391304	27304.179631
min	98001.000000	47.155900	-122.519000	399.000000	651.000000
25%	98033.000000	47.471000	-122.328000	1490.000000	5100.000000
50%	98065.000000	47.571800	-122.230000	1840.000000	7620.000000
75%	98118.000000	47.678000	-122.125000	2360.000000	10083.000000
max	98199.000000	47.777600	-121.315000	6210.000000	871200.000000

Missing Data

In this section we will explore the initial dataframe and understand missing values, performing data cleaning, wrangling and transformation of data where appropriate.

```
In [ ]: #Getting the number of missing values in each column
print(housing_data.isnull().sum())
```

```
id          0
date        0
price       0
bedrooms    1134
bathrooms   1068
sqft_living 1110
sqft_lot    1044
floors      0
waterfront  0
view        0
condition   0
grade       0
sqft_above  0
sqft_basement 0
yr_built    0
yr_renovated 0
zipcode     0
lat         0
long        0
sqft_living15 0
sqft_lot15  0
dtype: int64
```

Now I want to get a feel for how many 0's are in each column of the dataset

```
In [ ]: #Counting the number of zeros in each column
print(housing_data.isin([0]).sum())
```

```
id          0
date        0
price       0
bedrooms    11
bathrooms    8
sqft_living  0
sqft_lot     0
floors      0
waterfront   21450
view         19489
condition    0
grade       0
sqft_above  0
sqft_basement 13126
yr_built    0
yr_renovated 20699
zipcode     0
lat         0
long        0
sqft_living15 0
sqft_lot15  0
dtype: int64
```

At this point I am going to take advantage, and create two new dataframes, one will be cleaned and imputed data, the other will be simply removing all null values from the dataframe. We will start first with the simple removing all nulls values from the dataframe.

```
In [ ]: # Removing all NaN values from the original dataframe
housing_data_dropped_nan = housing_data.dropna()
print(housing_data_dropped_nan.shape)
```

```
(17618, 21)
```

At this point we grouped missing values for bedrooms and bathrooms by zipcode and imputed the most common bedroom or bathroom for that zipcode to fill in missing NaN values. By using the mode we get the most common

```
In [ ]: # Moving the original dataset to a cleaned dataset to avoid modifying the original dataset
cleaned_data = housing_data.copy()

# Handling the missing values in bedrooms column by grouping by zip code and using the mode of the bedrooms in each zip code
# Grouping by zip code and passing bedrooms into the transform function
cleaned_data['bedrooms'] = cleaned_data.groupby('zipcode')['bedrooms'].transform(lambda x: x.fillna(x.mode()[0]))
#count number of missing values in the bedrooms column
print(cleaned_data['bedrooms'].isnull().sum())

# Handling the missing values in the bathrooms column by grouping by zip code then using the mode of the bathrooms in each zip code
# Grouping by zip code and bedroom and passing bathrooms into the transform function
cleaned_data['bathrooms'] = cleaned_data.groupby(['zipcode'])['bathrooms'].transform(lambda x: x.fillna(x.mode()[0]))
#count number of missing values in the bathrooms column
print(cleaned_data['bathrooms'].isnull().sum())

0
0
```

Now for the missing values in the columns 'sqft_living' and 'sqft_lot' we will use the same groupby approach, leveraging the zip codes. However, for these values it would be acceptable to use the mean as opposed to the mode of the column.

```
In [ ]: #Handling the missing values from the 'sqft_living' column by grouping by zip code and using the mean of the 'sqft_living' in each zip code
# Grouping by zip code and passing sqft_living into the transform function
cleaned_data['sqft_living'] = cleaned_data.groupby('zipcode')['sqft_living'].transform(lambda x: x.fillna(x.mean()))
#count number of missing values in the sq_ft column
print(cleaned_data['sqft_living'].isnull().sum())

#Handling the missing values from the 'sqft_lot' column by grouping by zip code and using the mean of the 'sqft_lot' in each zip code
# Grouping by zip code and passing sqft_lot into the transform function
cleaned_data['sqft_lot'] = cleaned_data.groupby('zipcode')['sqft_lot'].transform(lambda x: x.fillna(x.mean()))
#count number of missing values in the sqft_lot column
print(cleaned_data['sqft_lot'].isnull().sum())

print(cleaned_data.isnull().sum())
print(cleaned_data.shape)

#Changing values greater than 1 in 'view' to 1
cleaned_data['view'] = cleaned_data['view'].apply(lambda x: 1 if x > 0 else x)

0
0
id            0
date          0
price         0
bedrooms      0
bathrooms     0
sqft_living   0
sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
zipcode       0
lat           0
long          0
sqft_living15 0
sqft_lot15    0
dtype: int64
(21613, 21)
```

Section 2: Data Analysis and Visualizations

Identify Categorical, Numerical, and Ordinal Data

Categorical data includes waterfront, yr_built, yr_renovated, and zip code.

Numerical data includes date, price, sqft_living, sqft_lot, sqft_above, sqft_basement, lat, long, sqft_living15, sqft_lot15.

Ordinal data includes bedroom, bathroom, floors, view, condition and grade.

Measures of Centrality

```
In [ ]: #printing the descriptive statistics of the cleaned and imputed dataframe
print(cleaned_data.describe())

#calculating mode and median for price
print("\nMode for price is ", cleaned_data['price'].mode()[0])
print("Median for price is ", cleaned_data['price'].median())
#calculating 25% and 75% quartile
pq3,pq1 = np.percentile(cleaned_data['price'], [75,25])
#calculating IQR
iqr = pq3-pq1
print("IQR for price is ", iqr)

#calculating mode and median for sqft_living
print("\nMode for sqft_living is ", cleaned_data['sqft_living'].mode()[0])
print("Median for sqft_living is ", cleaned_data['sqft_living'].median())
#calculating 25% and 75% quartile
sq3,sq1 = np.percentile(cleaned_data['sqft_living'], [75,25])
#calculating IQR
sq_iqr = sq3-sq1
print("IQR for sqft_living is ", sq_iqr)

#calculating mode and median for sqft_lot
print("\nMode for sqft_lot is ", cleaned_data['sqft_lot'].mode()[0])
print("Median for sqft_lot is ", cleaned_data['sqft_lot'].median())
#calculating 25% and 75% quartile
sq13,sq11 = np.percentile(cleaned_data['sqft_lot'], [75,25])
#calculating IQR
sq1_iqr = sq13-sq11
print("IQR for sqft_lot is ", sq1_iqr)

#calculating mode and median for sqft_above
```

```

print("\nMode for sqft_above is ", cleaned_data['sqft_above'].mode()[0])
print("Median for sqft_above is ", cleaned_data['sqft_above'].median())
#calculating 25% and 75% quartile
sq3,sq1 = np.percentile(cleaned_data['sqft_above'], [75,25])
#calculating IQR
sq_iqr = sq3-sq1
print("IQR for sqft_above is ", sq_iqr)

#calculating mode and median for sqft_basement
print("\nMode for sqft_basement is ", cleaned_data['sqft_basement'].mode()[0])
print("Median for sqft_basement is ", cleaned_data['sqft_basement'].median())
#calculating 25% and 75% quartile
sqb3,sqb1 = np.percentile(cleaned_data['sqft_basement'], [75,25])
sqb_iqr = sqb3-sqb1
print("IQR for sqft_basement is ", sqb_iqr)

```

	id	price	bedrooms	bathrooms	sqft_living \
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000
mean	4.580302e+09	5.400881e+05	3.360015	2.102438	2080.786053
std	2.876566e+09	3.671272e+05	0.910834	0.768081	895.764068
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.500000	1450.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1930.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2530.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	12050.000000

	sqft_lot	floors	waterfront	view	condition \
count	2.161300e+04	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.519622e+04	1.494309	0.007542	0.098274	3.409430
std	4.062404e+04	0.539989	0.086517	0.297692	0.650743
min	5.200000e+02	1.000000	0.000000	0.000000	1.000000
25%	5.100000e+03	1.000000	0.000000	0.000000	3.000000
50%	7.700000e+03	1.500000	0.000000	0.000000	3.000000
75%	1.089100e+04	2.000000	0.000000	0.000000	4.000000
max	1.651359e+06	3.500000	1.000000	1.000000	5.000000

	grade	sqft_above	sqft_basement	yr_built	yr_renovated \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	7.656873	1788.390691	291.509045	1971.005136	84.402258
std	1.175459	828.090978	442.575043	29.373411	401.679240
min	1.000000	290.000000	0.000000	1900.000000	0.000000
25%	7.000000	1190.000000	0.000000	1951.000000	0.000000
50%	7.000000	1560.000000	0.000000	1975.000000	0.000000
75%	8.000000	2210.000000	560.000000	1997.000000	0.000000
max	13.000000	9410.000000	4820.000000	2015.000000	2015.000000

	zipcode	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	98077.939805	47.560053	-122.213896	1986.552492	12768.455652
std	53.505026	0.138564	0.140828	685.391304	27304.179631
min	98001.000000	47.155900	-122.519000	399.000000	651.000000
25%	98033.000000	47.471000	-122.328000	1490.000000	5100.000000
50%	98065.000000	47.571800	-122.230000	1840.000000	7620.000000
75%	98118.000000	47.678000	-122.125000	2360.000000	10083.000000
max	98199.000000	47.777600	-121.315000	6210.000000	871200.000000

```

Mode for price is 350000.0
Median for price is 450000.0
IQR for price is 323050.0

```

```

Mode for sqft_living is 1300.0
Median for sqft_living is 1930.0
IQR for sqft_living is 1080.0

```

```

Mode for sqft_lot is 5000.0
Median for sqft_lot is 7700.0
IQR for sqft_lot is 5791.0

```

```

Mode for sqft_above is 1300
Median for sqft_above is 1560.0
IQR for sqft_above is 1020.0

```

```

Mode for sqft_basement is 0
Median for sqft_basement is 0.0
IQR for sqft_basement is 560.0

```

Distribution Visualizations

In this section, we will create visualizations for the distribution of data.

Here we are creating a histogram with a density plot for house price. Our data is positively skewed and most of the house price is between 321,950to645,000.

```

In [ ]: from matplotlib.pyplot import figure
figure(num=None, figsize=(15,7), dpi=256, facecolor='w', edgecolor='r')

#Creating a histogram with a density plot
p1 = sns.histplot(cleaned_data['price'], bins=100, kde=True, color = 'blue', edgecolor='red')

#Adding Labels and title
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.title("Histogram with Density Plot for House Price")
plt.ticklabel_format(style='plain')

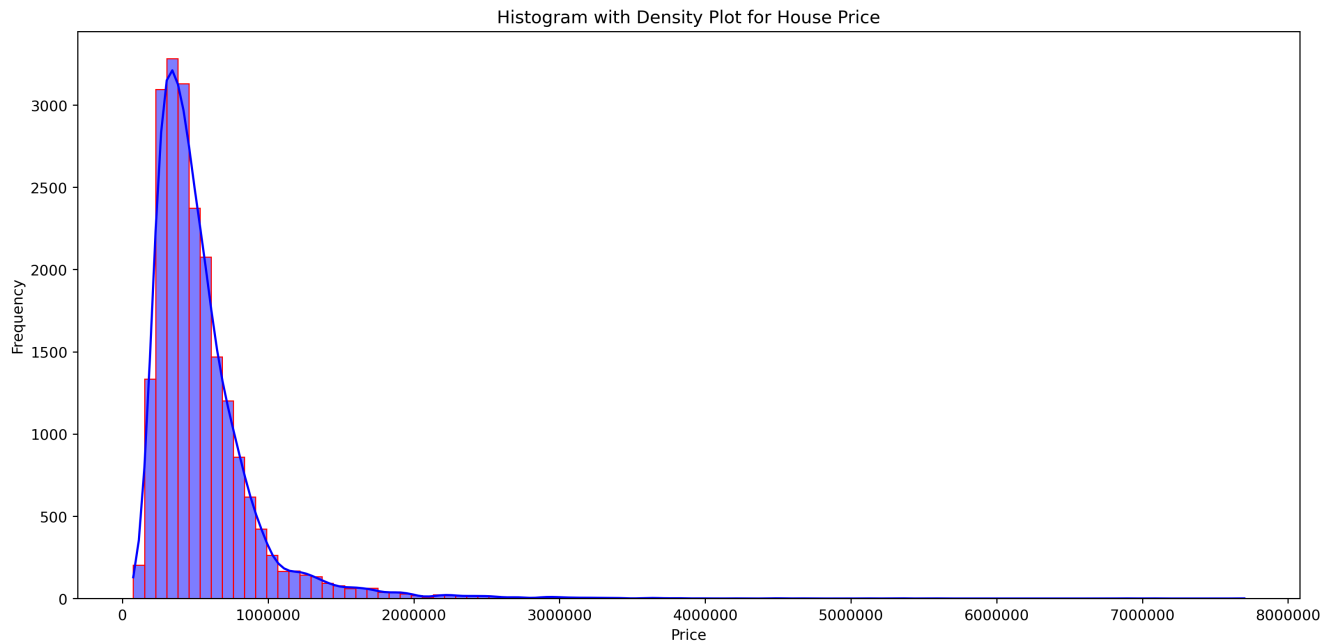
#showing the number of counts for each xLabel
#for i in p1.containers:
    #p1.bar_label(i,)

```

```

c:\Users\19noa\miniconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to Na
N before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```




We will now create a bar chart that shows the frequency for the number of bedrooms in a house.

```
In [ ]: plt.figure(figsize = (25,10))
# print(cleaned_data['bedrooms'].value_counts())

# Bar chart that shows the frequency for the number of bedrooms in a house
bed1 = sns.countplot(y='bedrooms', data = cleaned_data, color='skyblue', order = cleaned_data['bedrooms'].value_counts().index)

# scaling it with logarithmic scale to clearly show the bar plot
# bed1.set_xscale("log")
# ticks = [1, 10, 100, 1000, 10000, 13000]
# bed1.set_xticks(ticks)
# bed1.set_xticklabels(ticks)
plt.title("Bar chart with bedroom count")

# showing the number of counts for each xlabel
for i in bed1.containers:
    bed1.bar_label(i,)
```




We will now create a bar chart that shows the frequency for the number of bathrooms in a house.

```
In [ ]: plt.figure(figsize = (25,15))
# print(cleaned_data['bathrooms'].value_counts())

# Bar chart that shows the frequency for the number of bathrooms in a house.
bath1 = sns.countplot(y='bathrooms', data = cleaned_data, color='lightgreen', order = cleaned_data['bathrooms'].value_counts().index)

# scaling it with logarithmic scale to clearly show the bar plot
# bath1.set_xscale("log")
# ticks = [1, 10, 100, 1000, 10000]
# bath1.set_xticks(ticks)
# bath1.set_xticklabels(ticks)
plt.title("Bar chart with bathroom count")

# showing the number of counts for each xlabel
for i in bath1.containers:
    bath1.bar_label(i,)
```



We will now create a histogram with a density plot for sqft_living. The data is positively skewed and most of the sqft_living is between 290 square feet to 2530 square feet.


```
In [ ]: figure(num=None, figsize=(15,7), dpi=256, facecolor='w', edgecolor='r')

#Creating a histogram with a density plot
living1 = sns.histplot(cleaned_data['sqft_living'], bins=80, kde=True, color = 'blue', edgecolor='red')

#Adding Labels and title
plt.xlabel('sqft_living')
plt.ylabel('Frequency')
plt.title("Histogram with Density Plot for sqft_living")
plt.ticklabel_format(style='plain')

#showing the number of counts for each xlabel
#for i in living1.containers:
#    living1.bar_label(i,)
```

c:\Users\19noa\miniconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to Na N before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



We will now create a histogram with a density plot for sqft_lot. The data is positively skewed and most of the sqft_lot is between 5100 square feet to 10891 square feet.


```
In [ ]: figure(num=None, figsize=(20,18), dpi=256, facecolor='w', edgecolor='r')

#Creating a histogram with a density plot
lot1 = sns.histplot(cleaned_data['sqft_lot'], bins=99, kde=True, color = 'blue', edgecolor='red')

#Adding Labels and title
plt.xlabel('sqft_lot')
plt.ylabel('Frequency')
plt.title("Histogram with Density Plot for sqft_lot")
plt.ticklabel_format(style='plain')

#showing the number of counts for each xLabel
#for i in lot1.containers:
#    lot1.bar_label(i,)
```

c:\Users\19noa\miniconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to Na N before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



We will now create a histogram with a density plot for sqft_above. The data is positively skewed and most of the sqft_above is between 1190 square feet to 2210 square feet.

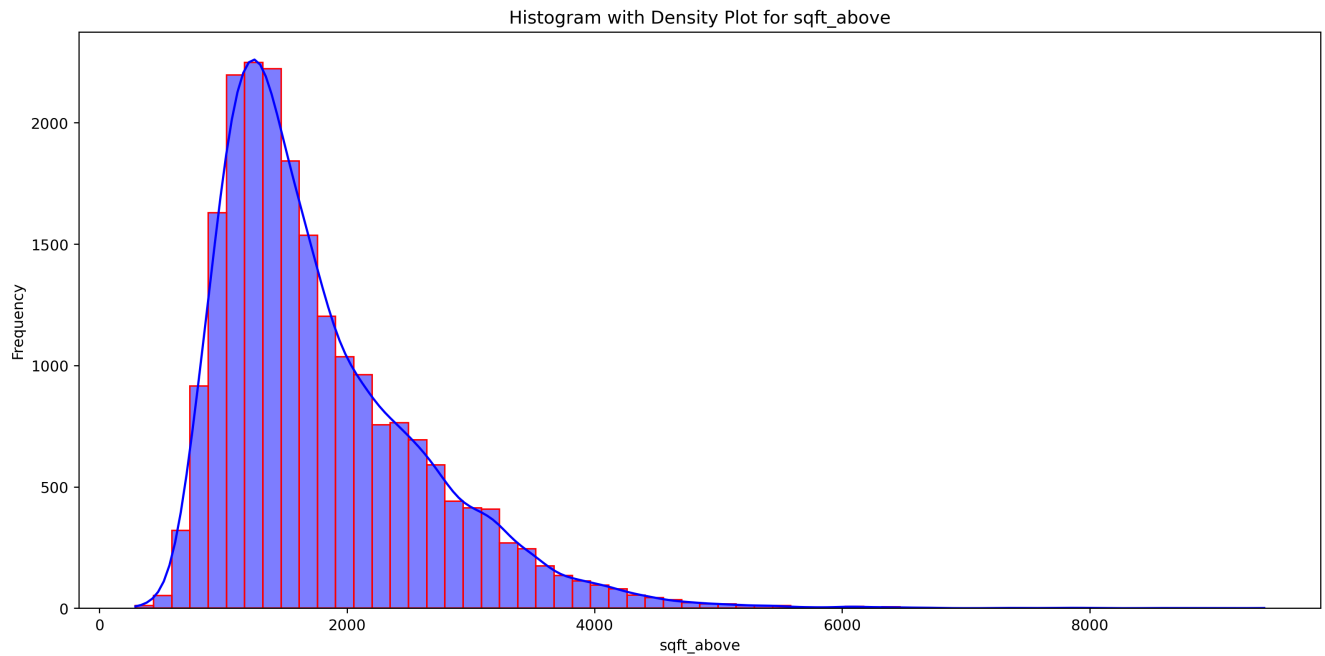
```
In [ ]: figure(num=None, figsize=(15,7), dpi=256, facecolor='w', edgecolor='r')

#Creating a histogram with a density plot
above1 = sns.histplot(cleaned_data['sqft_above'], bins=62, kde=True, color = 'blue', edgecolor='red')

#Adding Labels and title
plt.xlabel('sqft_above')
plt.ylabel('Frequency')
plt.title("Histogram with Density Plot for sqft_above")
plt.ticklabel_format(style='plain')

#showing the number of counts for each xLabel
#for i in above1.containers:
#    above1.bar_label(i,)
```

c:\Users\19noa\miniconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to Na N before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



We will now create a histogram with a density plot for sqft_basement. The data is positively skewed and most of the sqft_basement is between 0 square feet to 560 square feet.

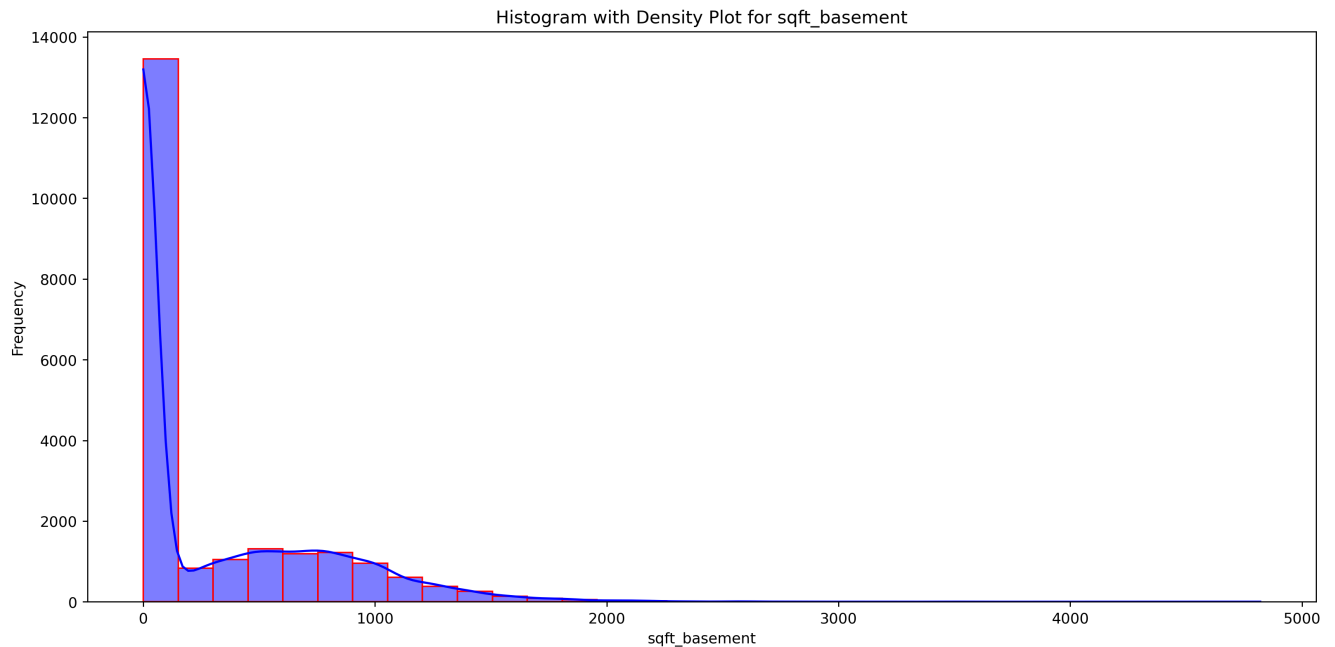
```
In [ ]: figure(num=None, figsize=(15,7), dpi=256, facecolor='w', edgecolor='r')

#Creating a histogram with a density plot
base1 = sns.histplot(cleaned_data['sqft_basement'], bins=32, kde=True, color = 'blue', edgecolor='red')

#Adding Labels and title
plt.xlabel('sqft_basement')
plt.ylabel('Frequency')
plt.title("Histogram with Density Plot for sqft_basement")
plt.ticklabel_format(style='plain')

#showing the number of counts for each xlabel
#for i in base1.containers:
#    base1.bar_label(i,)
```

c:\Users\19noa\miniconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to Na before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



We will now create a bar chart that shows the frequency for the number of floors in a house.

```
In [ ]: plt.figure(figsize = (40,5))


#Bar chart that shows the frequency for the number of floors in a house.
floors1 = sns.countplot(y='floors', data = cleaned_data, color='grey', order = cleaned_data['floors'].value_counts().index)

#scaling it with Logarithmic scale to clearly show the bar plot
#floors1.set_xscale("log")
#ticks = [1, 10, 100, 1000, 10000, 12000]
#floors1.set_xticks(ticks)
#floors1.set_xtickLabels(ticks)
```



```
plt.title("Bar chart with floors count")

#showing the number of counts for each xlabel
for i in floors1.containers:
    floors1.bar_label(i,)
```




We will now create a bar chart that shows the frequency of there being a waterfront in a house.

```
In [ ]: plt.figure(figsize = (25,5))
        #print(cleaned_data['waterfront'].value_counts())

        #Bar chart that shows the frequency of there being a waterfront in a house.
        water1 = sns.countplot(y='waterfront', data = cleaned_data, color='yellow', order = cleaned_data['waterfront'].value_counts().index)

        #scaling it with logarithmic scale to clearly show the bar plot
        #water1.set_xscale("log")
        #ticks = [1, 10, 100, 1000, 10000, 25000]
        #water1.set_xticks(ticks)
        #water1.set_xticklabels(ticks)
        plt.title("Bar chart for houses with waterfront count")

        #showing the number of counts for each xlabel
        for i in water1.containers:
            water1.bar_label(i,)
```




We will now create a bar chart that shows the frequency of there being a view in the house.

```
In [ ]: plt.figure(figsize = (30,5))
        #print(cleaned_data['view'].value_counts())

        #Bar chart that shows the frequency of there being a view in the house.
        view1 = sns.countplot(y='view', data = cleaned_data, color='violet', order = cleaned_data['view'].value_counts().index)

        #scaling it with logarithmic scale to clearly show the bar plot
        #view1.set_xscale("log")
        #ticks = [1, 10, 100, 1000, 10000, 25000]
        #view1.set_xticks(ticks)
        #view1.set_xticklabels(ticks)
        plt.title("Bar chart for houses with view count")

        #showing the number of counts for each xlabel
        for i in view1.containers:
            view1.bar_label(i,)
```




We will now create a bar chart that shows the frequency of the grade of the house.

```
In [ ]: plt.figure(figsize = (30,5))
        #print(cleaned_data['grade'].value_counts())

        #Bar that shows the frequency of the grade of the house.
        grade1 = sns.countplot(y='grade', data = cleaned_data, color='maroon', order = cleaned_data['grade'].value_counts().index)

        #scaling it with logarithmic scale to clearly show the bar plot
        #grade1.set_xscale("log")
        #ticks = [1, 10, 100, 1000, 10000]
        #grade1.set_xticks(ticks)
        #grade1.set_xticklabels(ticks)
        plt.title("Bar chart with grade frequency of the house")

        #showing the number of counts for each xlabel
        for i in grade1.containers:
            grade1.bar_label(i,)
```




Using a boxplot to identify outliers, there are approximately 1146 rows that contains an outlier for price and will be kept in the dataset.

```
In [ ]: price_boxplot = cleaned_data.boxplot(column=['price'])

# plot title
plt.title ('Boxplot of Price Data')
plt.show()

#calculating 25% and 75% quartile
price3, price1 = np.percentile(cleaned_data['price'], [75, 25])
price_iqr = price3 - price1
#calculating lower and upper bound
price_lower = cleaned_data['price'].min()
price_upper = price3 + 1.5 * price_iqr
print("Any value for price less than", price_lower, "dollars is an outlier.", "\nAny values greater than", price_upper, "dollars is an outlier.")
cleaned_data[cleaned_data.price > price_upper]
```



Any value for price less than 75000.0 dollars is an outlier
Any values greater than 1129575.0 dollars is an outlier.

```
Out [ ]: 
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	
	5	7237550310	20140512T000000	1225000.0	4.0	4.50	5420.000000	101930.000000	1.0	0	0	...	11	3890	1530	2001	0	98053 47
21	2524049179	20140826T000000	2000000.0	3.0	2.75	3050.000000	44867.000000	1.0	0	1	...	9	2330	720	1968	0	98040 47	
49	822039084	20150311T000000	1350000.0	3.0	2.50	2753.000000	65005.000000	1.0	1	1	...	9	2165	588	1953	0	98070 47	
69	1802000060	20140612T000000	1325000.0	5.0	2.25	3200.000000	12925.178451	1.0	0	0	...	8	1600	1600	1965	0	98004 47	
125	4389200955	20150302T000000	1450000.0	4.0	2.75	2750.000000	17789.000000	1.5	0	0	...	8	1980	770	1914	1992	98004 47	
...	
21568	524059330	20150130T000000	1700000.0	4.0	3.50	3830.000000	8963.000000	2.0	0	0	...	10	3120	710	2014	0	98004 47	
21576	9253900271	20150107T000000	3567000.0	5.0	4.50	4850.000000	10584.000000	2.0	1	1	...	10	3540	1310	2007	0	98008 47	
21590	7430200100	20140514T000000	1222500.0	4.0	3.50	2634.544153	9444.000000	1.5	0	0	...	11	3110	1800	2007	0	98074 47	
21597	191100405	20150421T000000	1575000.0	4.0	3.25	3410.000000	10125.000000	2.0	0	0	...	10	3410	0	2007	0	98040 47	
21600	249000205	20141015T000000	1537000.0	5.0	3.75	4470.000000	8088.000000	2.0	0	0	...	11	4470	0	2008	0	98004 47	

1146 rows × 21 columns

Using the boxplot, there are approximately 518 rows with outliers within the dataset. The extreme outlier of 33 bedrooms was analyzed and updated to be 3 bedrooms as there seems to be a typo. The sqft_living of 1620 and sqft_lot of 8049 with 1 floor does not accurately represent a 33 bedroom house. The 0 bedrooms were also removed.

```
In [ ]: bedrooms_boxplot = cleaned_data.boxplot(column=['bedrooms'])

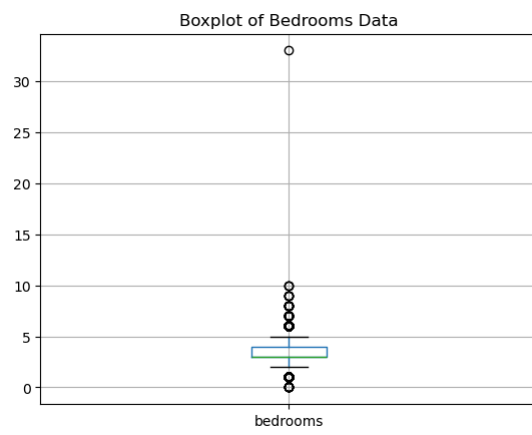
# plot title
plt.title ('Boxplot of Bedrooms Data')
plt.show()

#calculating 25% and 75% quartile
bed3, bed1 = np.percentile(cleaned_data['bedrooms'], [75, 25])
bed_iqr = bed3 - bed1
#calculating lower and upper bound
bed_lower = bed1 - 1.5 * bed_iqr
bed_upper = bed3 + 1.5 * bed_iqr
print("Any value for bedrooms less than", bed_lower, "bedrooms is an outlier.", "\nAny values greater than", bed_upper, "bedrooms is an outlier.")
print(cleaned_data[(cleaned_data.bedrooms < bed_lower) | (cleaned_data.bedrooms > bed_upper)].sort_values('bedrooms', ascending = True))

#replacing the 33 bedroom as a 3 bedroom
cleaned_data['bedrooms'].replace(33, 3, inplace=True)
print(cleaned_data[(cleaned_data.bedrooms < bed_lower) | (cleaned_data.bedrooms > bed_upper)].sort_values('bedrooms', ascending = True))

#removing all 0 bedrooms
cleaned_data = cleaned_data.loc[~(cleaned_data['bedrooms'] == 0)]
```

```
print(cleaned_data['bedrooms'].value_counts())
cleaned_data.sort_values('bedrooms', ascending = True)
```



Any value for bedrooms less than 1.5 bedrooms is an outlier
Any values greater than 5.5 bedrooms is an outlier.

	id	date	price	bedrooms	bathrooms	\
4868	6896300380	20141002T000000	228000.0	0.0	1.00	
18379	1222029077	20141029T000000	265000.0	0.0	0.75	
12653	7849202299	20150218T000000	320000.0	0.0	2.50	
14423	9543000205	20150413T000000	139950.0	0.0	2.50	
9773	3374500520	20150429T000000	355000.0	0.0	0.00	
...	
4235	2902200015	20150106T000000	700000.0	9.0	3.00	
13314	627300145	20140814T000000	1148000.0	10.0	5.25	
19254	8812401450	20141229T000000	660000.0	10.0	3.00	
15161	5566100170	20141029T000000	650000.0	10.0	2.00	
15870	2402100895	20140625T000000	640000.0	33.0	1.75	

	sqft_living	sqft_lot	floors	waterfront	view	...	grade	\
4868	390.0	5900.0	1.0	0	0	...	4	
18379	384.0	213444.0	1.0	0	0	...	4	
12653	1490.0	7111.0	2.0	0	0	...	7	
14423	844.0	4269.0	1.0	0	0	...	7	
9773	2460.0	8049.0	2.0	0	0	...	8	
...	
4235	3680.0	4400.0	2.0	0	0	...	7	
13314	4590.0	10920.0	1.0	0	1	...	9	
19254	2920.0	3745.0	2.0	0	0	...	7	
15161	3610.0	11914.0	2.0	0	0	...	7	
15870	1620.0	6000.0	1.0	0	0	...	7	

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	\
4868	390	0	1953	0	98118	47.5260	
18379	384	0	2003	0	98070	47.4177	
12653	1490	0	1999	0	98065	47.5261	
14423	844	0	1913	0	98001	47.2781	
9773	2460	0	1990	0	98031	47.4095	
...	
4235	2830	850	1908	0	98102	47.6374	
13314	2500	2090	2008	0	98004	47.5861	
19254	1860	1060	1913	0	98105	47.6635	
15161	3010	600	1958	0	98006	47.5705	
15870	1040	580	1947	0	98103	47.6878	

	long	sqft_living15	sqft_lot15
4868	-122.261	2170	6000
18379	-122.491	1920	224341
12653	-121.826	1500	4675
14423	-122.250	1380	9600
9773	-122.168	2520	8050
...
4235	-122.324	1960	2450
13314	-122.113	2730	10400
19254	-122.320	1810	3745
15161	-122.175	2040	11914
15870	-122.331	1330	4700

[518 rows x 21 columns]

	id	date	price	bedrooms	bathrooms	\
4868	6896300380	20141002T000000	228000.0	0.0	1.00	
18379	1222029077	20141029T000000	265000.0	0.0	0.75	
12653	7849202299	20150218T000000	320000.0	0.0	2.50	
14423	9543000205	20150413T000000	139950.0	0.0	2.50	
9773	3374500520	20150429T000000	355000.0	0.0	0.00	
...	
16844	8823900290	20150317T000000	1400000.0	9.0	4.00	
4235	2902200015	20150106T000000	700000.0	9.0	3.00	
13314	627300145	20140814T000000	1148000.0	10.0	5.25	
15161	5566100170	20141029T000000	650000.0	10.0	2.00	
19254	8812401450	20141229T000000	660000.0	10.0	3.00	

	sqft_living	sqft_lot	floors	waterfront	view	...	grade	\
4868	390.0	5900.0	1.0	0	0	...	4	
18379	384.0	213444.0	1.0	0	0	...	4	
12653	1490.0	7111.0	2.0	0	0	...	7	
14423	844.0	4269.0	1.0	0	0	...	7	
9773	2460.0	8049.0	2.0	0	0	...	8	
...	
16844	4620.0	5508.0	2.5	0	0	...	11	
4235	3680.0	4400.0	2.0	0	0	...	7	
13314	4590.0	10920.0	1.0	0	1	...	9	
15161	3610.0	11914.0	2.0	0	0	...	7	
19254	2920.0	3745.0	2.0	0	0	...	7	

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	\
4868	390	0	1953	0	98118	47.5260	
18379	384	0	2003	0	98070	47.4177	
12653	1490	0	1999	0	98065	47.5261	
14423	844	0	1913	0	98001	47.2781	
9773	2460	0	1990	0	98031	47.4095	
...	
16844	3870	750	1915	0	98105	47.6684	
4235	2830	850	1908	0	98102	47.6374	
13314	2500	2090	2008	0	98004	47.5861	
15161	3010	600	1958	0	98006	47.5705	
19254	1860	1060	1913	0	98105	47.6635	

	long	sqft_living15	sqft_lot15
4868	-122.261	2170	6000
18379	-122.491	1920	224341
12653	-121.826	1500	4675
14423	-122.250	1380	9600
9773	-122.168	2520	8050
...
16844	-122.309	2710	4320
4235	-122.324	1960	2450
13314	-122.113	2730	10400
15161	-122.175	2040	11914
19254	-122.320	1810	3745

[517 rows x 21 columns]

bedrooms
3.0 10275
4.0 6665
2.0 2617

```
5.0    1539
6.0     263
1.0     189
7.0      34
8.0      12
9.0       5
10.0      3
```

Name: count, dtype: int64

C:\Users\19noa\AppData\Local\Temp\ipykernel_52056\2335008484.py:17: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
cleaned_data['bedrooms'].replace(33, 3, inplace=True)
```

```
Out [ ]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	
	860	1723049033	20140620T000000	245000.0	1.0	0.75	380.0	15000.0000	1.0	0	0	...	5	380	0	1963	0	98168	47.4810
	14366	3333002450	20140708T000000	165000.0	1.0	1.00	850.0	8050.0000	1.0	0	0	...	6	850	0	1906	0	98118	47.5427
	18052	1352300580	20141114T000000	247000.0	1.0	1.00	460.0	10225.6875	1.0	0	0	...	4	460	0	1937	0	98055	47.4868
	21240	7174800094	20150420T000000	525000.0	1.0	1.50	1030.0	5923.0000	1.0	0	0	...	8	1030	0	1940	0	98105	47.6653
	18059	1773101530	20141218T000000	275000.0	1.0	1.00	520.0	4800.0000	1.0	0	0	...	5	520	0	1930	0	98106	47.5533

	6079	9822700190	20140808T000000	1280000.0	9.0	4.50	3650.0	5000.0000	2.0	0	0	...	8	2530	1120	1915	2010	98105	47.6604
	18443	8823901445	20150313T000000	934000.0	9.0	1.75	2820.0	4480.0000	2.0	0	0	...	7	1880	940	1918	0	98105	47.6654
	19254	8812401450	20141229T000000	660000.0	10.0	3.00	2920.0	3745.0000	2.0	0	0	...	7	1860	1060	1913	0	98105	47.6635
	13314	627300145	20140814T000000	1148000.0	10.0	5.25	4590.0	10920.0000	1.0	0	1	...	9	2500	2090	2008	0	98004	47.5861
	15161	5566100170	20141029T000000	650000.0	10.0	2.00	3610.0	11914.0000	2.0	0	0	...	7	3010	600	1958	0	98006	47.5705

21602 rows × 21 columns



Using the boxplot, there are approximately 252 rows with outliers within the dataset. The 0 bathrooms were removed from the dataset.

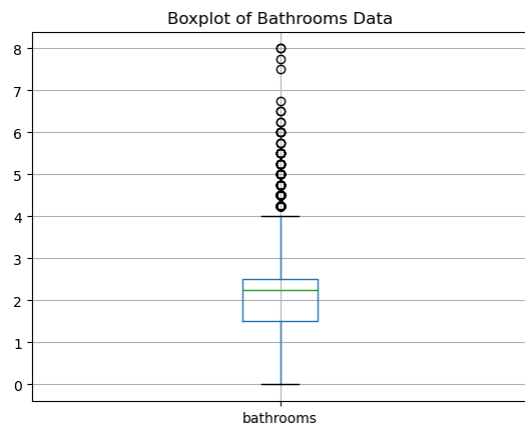
```
In [ ]:
```

```
bathrooms_boxplot = cleaned_data.boxplot(column=['bathrooms'])

# plot title
plt.title('Boxplot of Bathrooms Data')
plt.show()

#calculating 25% and 75% quartile
bath3,bath1 = np.percentile(cleaned_data['bathrooms'], [75,25])
bath_iqr = bath3-bath1
#calculating lower and upper bound
bath_lower = bath1 - 1.5*bath_iqr
bath_upper = bath3 + 1.5*bath_iqr
print("Any value for bathrooms less than",bath_lower, "bathrooms is an outlier.", "\nAny values greater than", bath_upper, "bathrooms is an outlier.")
print(cleaned_data[(cleaned_data.bathrooms < bath_lower) | (cleaned_data.bathrooms > bath_upper)].sort_values('bathrooms', ascending = True))

#removing all 0 bathrooms
cleaned_data = cleaned_data.loc[~(cleaned_data['bathrooms'] == 0)]
print(cleaned_data['bathrooms'].value_counts())
cleaned_data.sort_values('bathrooms', ascending = True)
```



Any value for bathrooms less than 0.0 bathrooms is an outlier.
Any values greater than 4.0 bathrooms is an outlier.

	id	date	price	bedrooms	bathrooms	\
11685	1126069045	20140620T000000	1135000.0	6.0	4.25	
7710	644200040	20140515T000000	1000000.0	5.0	4.25	
15022	2210500010	20140930T000000	2450000.0	7.0	4.25	
7280	922059169	20141201T000000	800000.0	6.0	4.25	
7236	1245002391	20141022T000000	1400000.0	5.0	4.25	
...
8092	1924059029	20140617T000000	4668000.0	5.0	6.75	
8546	424049043	20140811T000000	450000.0	9.0	7.50	
9254	9208900037	20140919T000000	6885000.0	6.0	7.75	
12777	1225069038	20140505T000000	2280000.0	3.0	8.00	
7252	6762700020	20141013T000000	7700000.0	6.0	8.00	

	sqft_living	sqft_lot	floors	waterfront	view	...	grade	\
11685	6900.000000	244716.0	2.0	0	0	...	9	
7710	3920.000000	16258.0	2.0	0	0	...	9	
15022	4670.000000	23115.0	2.0	0	1	...	11	
7280	5480.000000	189050.0	2.0	0	0	...	10	
7236	4230.000000	6907.0	2.0	0	0	...	10	
...
8092	9640.000000	13068.0	1.0	1	1	...	12	
8546	4050.000000	6504.0	2.0	0	0	...	7	
9254	3685.520833	31374.0	2.0	0	1	...	13	
12777	2584.620053	307752.0	3.0	0	1	...	12	
7252	12050.000000	27600.0	2.5	0	1	...	13	

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	\
11685	4820	2080	2002	0	98077	47.7506	
7710	2900	1020	1990	0	98004	47.5871	
15022	4670	0	1992	0	98039	47.6183	
7280	5140	340	1991	0	98031	47.4120	
7236	3450	780	2008	0	98033	47.6866	
...
8092	4820	4820	1983	2009	98040	47.5570	
8546	4050	0	1996	0	98144	47.5923	
9254	8860	1030	2001	0	98039	47.6305	
12777	9410	4130	1999	0	98053	47.6675	
7252	8570	3480	1910	1987	98102	47.6298	

	long	sqft_living15	sqft_lot15
11685	-122.012	4170	266587
7710	-122.192	2540	12131
15022	-122.227	3240	13912
7280	-122.168	2470	10429
7236	-122.205	2650	8076
...
8092	-122.210	3270	10454
8546	-122.301	1448	3866
9254	-122.240	4540	42730
12777	-121.986	4850	217800
7252	-122.323	3940	8800

[252 rows x 21 columns]

bathrooms

2.50	5713
1.00	4077
1.75	2942
2.25	1964
2.00	1828
1.50	1383
2.75	1121
3.00	713
3.50	690
3.25	558
3.75	150
4.00	128
4.50	94
4.25	77
0.75	67
4.75	23
5.00	21
5.25	10
5.50	10
1.25	8
6.00	5
0.50	4
5.75	3
8.00	2
6.25	2
6.50	2
6.75	1
7.50	1
7.75	1

Name: count, dtype: int64

Out [] :

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat
10424	7129800036	20150114T000000	109000.0	2.0	0.50	580.000000	6900.0	1.0	0	0	...	5	580	0	1941	0	98118	47.5135
12041	2991000160	20141212T000000	312500.0	4.0	0.50	2300.000000	5570.0	2.0	0	0	...	8	2300	0	1996	0	98092	47.3285
11674	7987400316	20140814T000000	255000.0	1.0	0.50	880.000000	1642.0	1.0	0	0	...	6	500	380	1910	0	98126	47.5732
2261	3971701455	20141003T000000	273000.0	2.0	0.50	1180.000000	7750.0	1.0	0	0	...	6	590	590	1945	0	98155	47.7690
21612	1523300157	20141015T000000	325000.0	2.0	0.75	1020.000000	1076.0	2.0	0	0	...	7	1020	0	2008	0	98144	47.5941
...
8092	1924059029	20140617T000000	4668000.0	5.0	6.75	9640.000000	13068.0	1.0	1	1	...	12	4820	4820	1983	2009	98040	47.5570
8546	424049043	20140811T000000	450000.0	9.0	7.50	4050.000000	6504.0	2.0	0	0	...	7	4050	0	1996	0	98144	47.5923
9254	9208900037	20140919T000000	6885000.0	6.0	7.75	3685.520833	31374.0	2.0	0	1	...	13	8860	1030	2001	0	98039	47.6305
12777	1225069038	20140505T000000	2280000.0	3.0	8.00	2584.620053	307752.0	3.0	0	1	...	12	9410	4130	1999	0	98053	47.6675
7252	6762700020	20141013T000000	7700000.0	6.0	8.00	12050.000000	27600.0	2.5	0	1	...	13	8570	3480	1910	1987	98102	47.6298

21598 rows x 21 columns

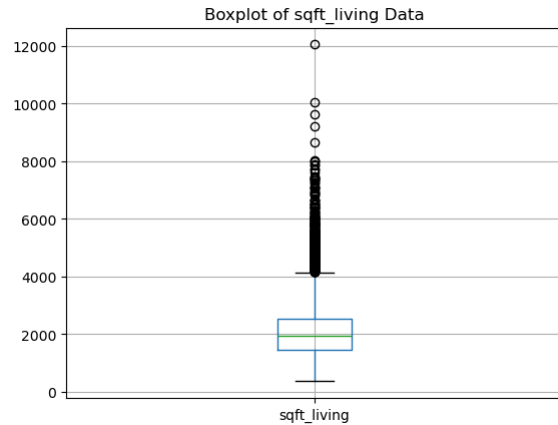


Using the boxplot, there are approximately 601 rows with outliers within the sqft_living column and will be kept in the dataset.

```
In [ ]: sqft_living_boxplot = cleaned_data.boxplot(column=['sqft_living'])

# plot title
plt.title ('Boxplot of sqft_living Data')
plt.show()

#calculating 25% and 75% quartile
sqft3,sqft1 = np.percentile(cleaned_data['sqft_living'], [75,25])
sqft_iqr = sqft3-sqft1
#calculating lower and upper bound
sqft_lower = cleaned_data['sqft_living'].min()
sqft_upper = sqft3 + 1.5*sqft_iqr
print("Any values less than",sqft_lower, "square foot is an outlier.", "\nAny values greater than",sqft_upper,"square foot is an outlier.")
cleaned_data[(cleaned_data.sqft_living > sqft_upper)].sort_values('sqft_living', ascending = False)
```



Any values less than 370.0 square foot is an outlier.
Any values greater than 4150.0 square foot is an outlier.

```
Out [ ]: 
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	I
7252	6762700020	20141013T000000	7700000.0	6.0	8.00	12050.0	27600.000000	2.5	0	1	...	13	8570	3480	1910	1987	98102	47.62
3914	9808700762	20140611T000000	7062500.0	5.0	4.50	10040.0	37325.000000	2.0	1	1	...	11	7680	2360	1940	2001	98004	47.65
8092	1924059029	20140617T000000	4668000.0	5.0	6.75	9640.0	13068.000000	1.0	1	1	...	12	4820	4820	1983	2009	98040	47.55
4411	2470100110	20140804T000000	5570000.0	5.0	5.75	9200.0	16813.145833	2.0	0	0	...	13	6200	3000	2001	0	98039	47.62
14556	2303900035	20140611T000000	2888000.0	5.0	6.25	8670.0	64033.000000	2.0	0	1	...	13	6120	2550	1965	2003	98177	47.72
...
11947	3303980140	20150402T000000	1150000.0	4.0	3.00	4160.0	13170.000000	2.0	0	0	...	11	3040	1120	2001	0	98059	47.51
11233	2655500241	20140814T000000	1699000.0	3.0	3.25	4160.0	35153.000000	3.0	0	1	...	12	3690	470	2001	0	98040	47.57
5961	8155800050	20150422T000000	1110000.0	3.0	4.00	4160.0	31796.000000	2.0	0	0	...	11	4160	0	1989	0	98053	47.66
19681	1266200140	20150506T000000	1850000.0	4.0	3.25	4160.0	10335.000000	2.0	0	0	...	10	4160	0	2014	0	98004	47.62
13448	2426059124	20141216T000000	1045000.0	4.0	3.25	4160.0	47480.000000	2.0	0	0	...	10	4160	0	1995	0	98072	47.72

601 rows × 21 columns

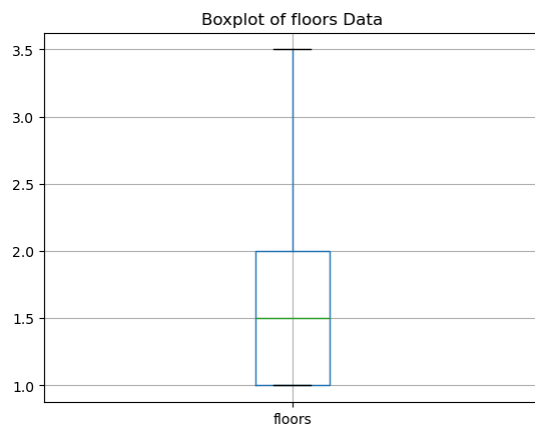
◀ ▶

Using the boxplot, there are no outliers in the floors column.

```
In [ ]: floors_boxplot = cleaned_data.boxplot(column=['floors'])

# plot title
plt.title ('Boxplot of floors Data')
plt.show()

#calculating 25% and 75% quartile
floors3,floors1 = np.percentile(cleaned_data['floors'], [75,25])
floors_iqr = floors3-floors1
#calculating lower and upper bound
floors_lower = cleaned_data['floors'].min()
floors_upper = floors3 + 1.5*floors_iqr
print("Any value less than",floors_lower, "floors is an outlier.", "\nAny values greater than",floors_upper, "floors is an outlier.")
cleaned_data[(cleaned_data.floors > floors_upper)].sort_values('floors', ascending = False)
```



Any value less than 1.0 floors is an outlier.
Any values greater than 3.5 floors is an outlier.

Out []: id date price bedrooms bathrooms sqft_living sqft_lot floors waterfront view ... grade sqft_above sqft_basement yr_built yr_renovated zipcode lat long sqft_living15 sqft_lot15

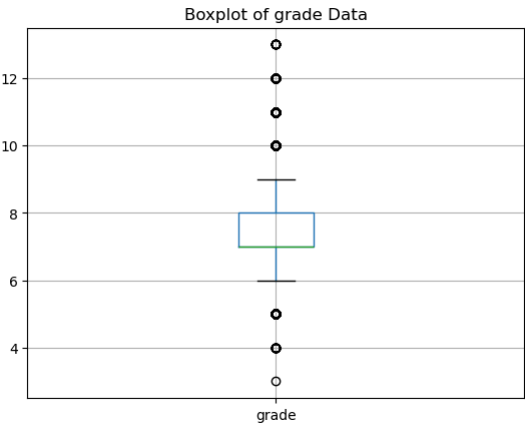
0 rows × 21 columns

Using the boxplot, there are approximately 1905 rows with outliers within the grade column and will be kept from the dataset.

```
In [ ]: grade_boxplot = cleaned_data.boxplot(column=['grade'])

# plot title
plt.title ('Boxplot of grade Data')
plt.show()

#calculating 25% and 75% quartile
grade3,grade1 = np.percentile(cleaned_data['grade'], [75,25])
grade_iqr = grade3-grade1
#calculating lower and upper bound
grade_lower = grade1 - 1.5*grade_iqr
grade_upper = grade3 + 1.5*grade_iqr
print("Any values less than",grade_lower, "for grade is an outlier.", "\nAny values greater than",grade_upper, "for grade is an outlier.")
cleaned_data[(cleaned_data.grade < grade_lower) | (cleaned_data.grade > grade_upper)].sort_values('grade', ascending = False)
```



Any values less than 5.5 for grade is an outlier.
Any values greater than 9.5 for grade is an outlier.

Out []:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode
6041	1725059316	20141120T000000	2385000.0	4.0	4.00	2382.779376	13296.000000	2.0	0	1	...	13	4900	1430	2000	0	98033 47
5451	7237501190	20141010T000000	1780000.0	4.0	3.25	4890.000000	13402.000000	2.0	0	0	...	13	4890	0	2004	0	98059 47
7252	6762700020	20141013T000000	7700000.0	6.0	8.00	12050.000000	27600.000000	2.5	0	1	...	13	8570	3480	1910	1987	98102 47
4411	2470100110	20140804T000000	5570000.0	5.0	5.75	9200.000000	16813.145833	2.0	0	0	...	13	6200	3000	2001	0	98039 47
19017	2303900100	20140911T000000	3800000.0	3.0	4.25	5510.000000	35000.000000	2.0	0	1	...	13	4910	600	1997	0	98177 47
...
7974	3760500240	20150512T000000	435000.0	2.0	0.75	750.000000	16321.000000	1.0	0	1	...	4	750	0	1936	0	98034 47
465	8658300340	20140523T000000	80000.0	1.0	0.75	430.000000	5050.000000	1.0	0	0	...	4	430	0	1912	0	98014 47
16340	6146600170	20140703T000000	100000.0	2.0	0.75	660.000000	5240.000000	1.0	0	0	...	4	660	0	1912	0	98032 47
7973	3122069029	20140619T000000	120000.0	2.0	1.00	990.000000	39964.000000	1.0	0	0	...	4	990	0	1945	0	98042 47
3223	2420069251	20150225T000000	262000.0	1.0	0.75	1841.699115	12981.000000	1.0	0	0	...	3	520	0	1920	0	98022 47

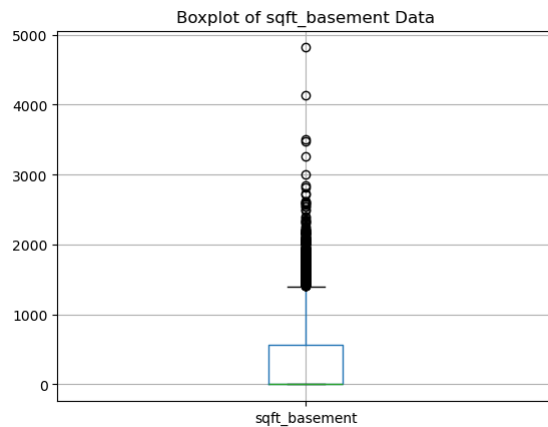
1905 rows × 21 columns

Using the boxplot, there are approximately 496 rows with outliers within the sqft_basement column and will be kept in the dataset.

```
In [ ]: sqft_base_boxplot = cleaned_data.boxplot(column=['sqft_basement'])

# plot title
plt.title ('Boxplot of sqft_basement Data')
plt.show()

#calculating 25% and 75% quartile
sqft_base3,sqft_base1 = np.percentile(cleaned_data['sqft_basement'], [75,25])
sqft_base_iqr = sqft_base3-sqft_base1
#calculating lower and upper bound
sqft_base_lower = cleaned_data['sqft_basement'].min()
sqft_base_upper = sqft_base3 + 1.5*sqft_base_iqr
print("Any values less than", sqft_base_lower, "square foot is an outlier", "\nAny values greater than", sqft_base_upper, "square foot is an outlier.")
cleaned_data[(cleaned_data.sqft_basement < sqft_base_lower) | (cleaned_data.sqft_basement > sqft_base_upper)].sort_values('sqft_basement', ascending = False)
```

Any values less than 0 square foot is an outlier
Any values greater than 1400.0 square foot is an outlier.

Out[]:		id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat
	8092	1924059029	20140617T000000	4668000.0	5.0	6.75	9640.000000	13068.0	1.0	1	1	...	12	4820	4820	1983	2009	98040	47.5570
	12777	1225069038	20140505T000000	2280000.0	3.0	8.00	2584.620053	307752.0	3.0	0	1	...	12	9410	4130	1999	0	98053	47.6675
	15482	624069108	20140812T000000	3200000.0	4.0	3.25	7000.000000	28206.0	1.0	1	1	...	12	3500	3500	1991	0	98075	47.5928
	7252	6762700020	20141013T000000	7700000.0	6.0	8.00	12050.000000	27600.0	2.5	0	1	...	13	8570	3480	1910	1987	98102	47.6298
	10085	7767000060	20140912T000000	1900000.0	5.0	4.25	6510.000000	16471.0	2.0	0	1	...	11	3250	3260	1980	0	98040	47.5758

	2209	269000240	20141030T000000	1050000.0	5.0	2.25	2168.729642	7680.0	1.0	0	1	...	8	1550	1410	1958	0	98199	47.6456
	4827	7366100080	20140731T000000	318000.0	5.0	2.50	2820.000000	9956.0	1.0	0	0	...	7	1410	1410	1967	0	98168	47.4715
	8976	1126049095	20140926T000000	450000.0	3.0	2.50	2820.000000	10208.0	1.0	0	1	...	8	1410	1410	1954	0	98028	47.7609
	4336	7738500475	20141212T000000	485000.0	3.0	3.25	2820.000000	6611.0	1.0	0	0	...	7	1410	1410	1958	0	98155	47.7473
	1539	3425059222	20141124T000000	1300000.0	6.0	3.50	2640.696203	32670.0	2.0	0	0	...	10	5153	1410	2002	0	98005	47.6078

496 rows × 21 columns



Correlation Matrix

This correlation matrix below uses the correlation function to show the correlation between all the variables. The heat map displays higher correlated items.

```
In [ ]: import seaborn as sb
columns_to_plot = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'grade']

chosen_data = cleaned_data[columns_to_plot]

correlation = chosen_data.corr()
corr_plot = sb.heatmap(correlation, annot=True)
print('Correlation Matrix: ')
print(correlation)
```

```
Correlation Matrix:
           price  bedrooms  bathrooms  sqft_living  sqft_lot  floors \
price      1.000000  0.315289  0.511176  0.688973  0.084613  0.256815
bedrooms   0.315289  1.000000  0.499751  0.572469  0.029216  0.180476
bathrooms  0.511176  0.499751  1.000000  0.717912  0.085412  0.480464
sqft_living 0.688973  0.572469  0.717912  1.000000  0.163590  0.346291
sqft_lot   0.084613  0.029216  0.085412  0.163590  1.000000 -0.004534
floors     0.256815  0.180476  0.480464  0.346291 -0.004534  1.000000
grade      0.667921  0.362937  0.646311  0.751033  0.109960  0.458806

           grade
price      0.667921
bedrooms   0.362937
bathrooms  0.646311
sqft_living 0.751033
sqft_lot   0.109960
floors     0.458806
grade      1.000000
```



Identifying the Independent and Dependent Variables


The independent variables include bedrooms, bathrooms, sqft_living, sqft_lot, floors, waterfront, view, and grade.

The dependent variable is price.

Exploratory Analysis via Visualizations

The below scatter plot displays the relationship between price and the number of bedrooms

```
In [ ]: #Price and Bedrooms
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'bedrooms', y = 'price', data=cleaned_data)
plt.title('Scatter plot for bedrooms and price')
plt.xlabel('bedrooms')
plt.ylabel('price')
plt.show()
```



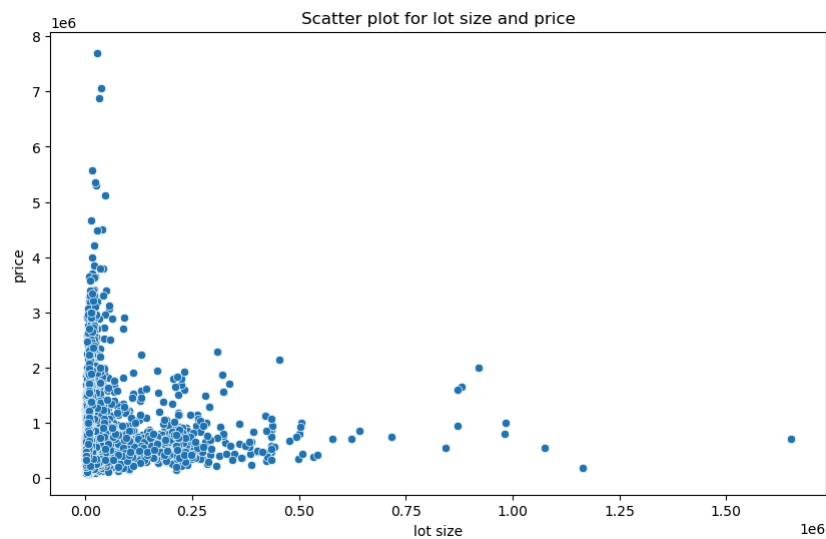
This scatter plot shows Price and Bathrooms

```
In [ ]: #Scatter for price and bathrooms
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'bathrooms', y = 'price', data=cleaned_data)
plt.title('Scatter plot for bathrooms and price')
plt.xlabel('bathrooms')
plt.ylabel('price')
plt.show()
```



This plot shows how the price and lot size are related. Oddly they don't seem to be positively correlated.

```
In [ ]: #Scatter for price and lot size
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'sqft_lot', y = 'price', data=cleaned_data)
plt.title('Scatter plot for lot size and price')
plt.xlabel('lot size')
plt.ylabel('price')
plt.show()
```



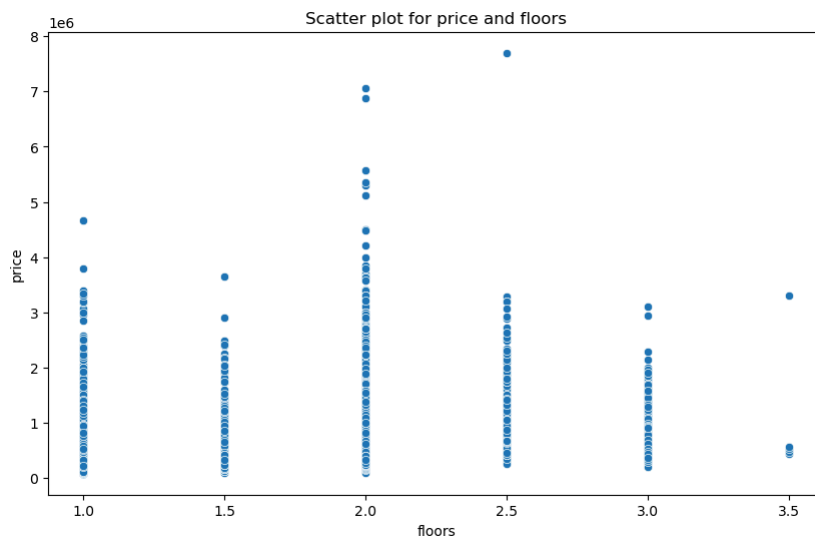
This plot shows the relationship between price and house square footage. This clearly shows a positive correlation between the square feet of the house and the price.

```
In [ ]: #Scatter for price and square foot living
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'sqft_living', y = 'price', data=cleaned_data)
plt.title('Scatter plot for square foot living and price')
plt.xlabel('square foot living')
plt.ylabel('price')
plt.show()
```



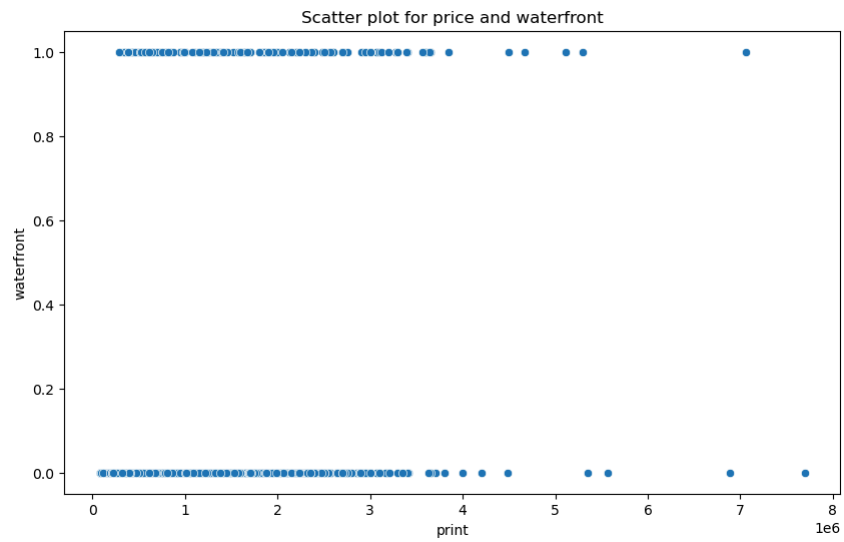
This scatter plot shows how floors and price are related. Strangely it looks like just because there are more floors the price doesn't necessarily increase. We see the majority of the higher priced houses are in the 2 floor category.

```
In [ ]: #Price and Floors EA scatter
plt.figure(figsize=(10,6))
sns.scatterplot(x='floors', y='price', data=cleaned_data)
plt.title('Scatter plot for price and floors')
plt.xlabel('floors')
plt.ylabel('price')
plt.show()
```



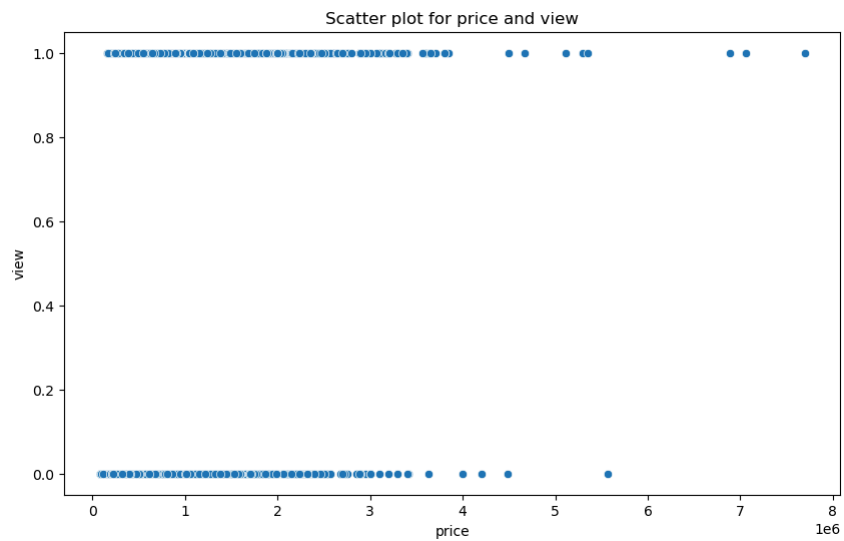
Here is the relationship between price and whether the house is waterfront or not and again these results were surprising. It doesn't show that being waterfront increases the price of the house.

```
In [ ]: #Scatter for price and waterfront
plt.figure(figsize=(10,6))
sns.scatterplot(x='price', y='waterfront', data=cleaned_data)
plt.title('Scatter plot for price and waterfront')
plt.xlabel('price')
plt.ylabel('waterfront')
plt.show()
```



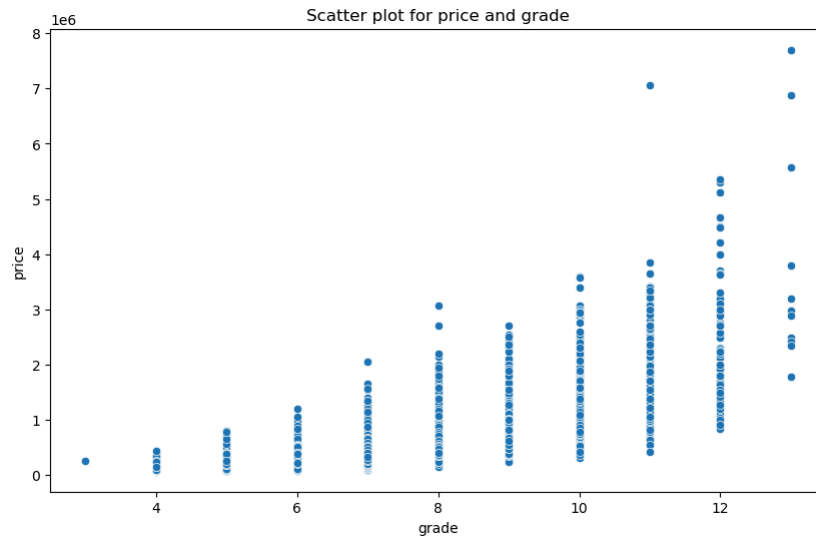
This shows price and view's relationship and it does seem that having a good view increases the value which is what we would probably assume.

```
In [ ]: #Scatter for price and view
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'price', y = 'view', data=cleaned_data)
plt.title('Scatter plot for price and view')
plt.xlabel('price')
plt.ylabel('view')
plt.show()
```



This below plot shows the relationship between price and grade which as expected appears to have a strong positive correlation

```
In [ ]: #Scatter for price and grade
plt.figure(figsize=(10,6))
sns.scatterplot(x = 'grade', y = 'price', data=cleaned_data)
plt.title('Scatter plot for price and grade')
plt.xlabel('grade')
plt.ylabel('price')
plt.show()
```



Section 3: Data Analytics - test

Supervised vs Unsupervised Learning

This is the regression analysis for sqft lot and sqft living against the price. R-squared value shows about 48% of the variability in the price depends on sqft living and lot. F-statistic is 9793 so with the number of observations and degrees of freedom we can determine this regression is most likely not statistically significant.

```
In [ ]: #regression for sqft_lot and sqft_living against price
```

```
import statsmodels.api as sm
```

```
In [ ]: X = cleaned_data[['sqft_lot', 'sqft_living']]
        Y = cleaned_data['price']
```

```
X = sm.add_constant(X)
```

```
linear_model = sm.OLS(Y, X)
```

```
results = linear_model.fit()
```

```
print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.475
Model:                  OLS      Adj. R-squared:             0.475
Method:                 Least Squares    F-statistic:         9789.
Date:                   Sat, 13 Apr 2024    Prob (F-statistic):    0.00
Time:                   19:27:27           Log-Likelihood:      -3.0042e+05
No. Observations:       21598             AIC:                6.009e+05
DF Residuals:           21595             BIC:                6.009e+05
DF Model:                2
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-4.774e+04	4577.872	-10.429	0.000	-5.67e+04	-3.88e+04
sqft_lot	-0.2609	0.045	-5.779	0.000	-0.349	-0.172
sqft_living	284.4136	2.048	138.861	0.000	280.399	288.428

```

=====
Omnibus:                 16351.480    Durbin-Watson:          1.975
Prob(Omnibus):            0.000    Jarque-Bera (JB):       914606.580
Skew:                     3.134    Prob(JB):               0.00
Kurtosis:                 34.258    Cond. No.               1.10e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.1e+05. This might indicate that there are strong multicollinearity or other numerical problems.

This is for the relation ship between bedrooms and price,R-squared is .099 and Fstatistic is 2387.

```
In [ ]: #Regression Model for price and bedrooms
```

```
X = cleaned_data[['bedrooms']]
Y = cleaned_data['price']
```

```
X = sm.add_constant(X)
```

```
linear_model = sm.OLS(Y, X)
```

```
results = linear_model.fit()
```

```
print(results.summary())
```

```

OLS Regression Results
=====
Dep. Variable:      price    R-squared:      0.099
Model:              OLS     Adj. R-squared:      0.099
Method:             Least Squares    F-statistic:      2384.
Date:               Sat, 13 Apr 2024    Prob (F-statistic):      0.00
Time:              19:27:27    Log-Likelihood:      -3.0626e+05
No. Observations:  21598    AIC:      6.125e+05
Df Residuals:      21596    BIC:      6.125e+05
Df Model:          1
Covariance Type:    nonrobust
=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
const      1.007e+05    9309.337     10.814     0.000     8.24e+04     1.19e+05
bedrooms   1.308e+05    2678.819     48.824     0.000     1.26e+05     1.36e+05
=====
Omnibus:      18899.226    Durbin-Watson:      1.962
Prob(Omnibus):      0.000    Jarque-Bera (JB):      1162326.629
Skew:         3.936    Prob(JB):      0.00
Kurtosis:     38.066    Cond. No.      14.7
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

This shows the relationship between Bathrooms and Price and we see an R-squared of .261 and F statistic is 7642

```
In [ ]: #Supervised regression for bathroom and price
```

```

X = cleaned_data[['bathrooms']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())

```

```

OLS Regression Results
=====
Dep. Variable:      price    R-squared:      0.261
Model:              OLS     Adj. R-squared:      0.261
Method:             Least Squares    F-statistic:      7639.
Date:               Sat, 13 Apr 2024    Prob (F-statistic):      0.00
Time:              19:27:27    Log-Likelihood:      -3.0412e+05
No. Observations:  21598    AIC:      6.082e+05
Df Residuals:      21596    BIC:      6.083e+05
Df Model:          1
Covariance Type:    nonrobust
=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
const      2.566e+04    6266.320     4.095     0.000     1.34e+04     3.79e+04
bathrooms   2.446e+05    2798.923    87.402     0.000     2.39e+05     2.5e+05
=====
Omnibus:      17317.711    Durbin-Watson:      1.956
Prob(Omnibus):      0.000    Jarque-Bera (JB):      879539.354
Skew:         3.476    Prob(JB):      0.00
Kurtosis:     33.480    Cond. No.      7.71
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Here is the regression for floors and prices R-squared is .066 and F-statistic is 1525.

```
In [ ]: #Supervised regression for price and floors
```

```

X = cleaned_data[['floors']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())

```

```

OLS Regression Results
=====
Dep. Variable:      price    R-squared:      0.066
Model:              OLS     Adj. R-squared:      0.066
Method:             Least Squares    F-statistic:      1525.
Date:               Sat, 13 Apr 2024    Prob (F-statistic):      2.27e-322
Time:              19:27:27    Log-Likelihood:      -3.0665e+05
No. Observations:  21598    AIC:      6.133e+05
Df Residuals:      21596    BIC:      6.133e+05
Df Model:          1
Covariance Type:    nonrobust
=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
const      2.792e+05    7106.838     39.284     0.000     2.65e+05     2.93e+05
floors     1.747e+05    4473.784     39.050     0.000     1.66e+05     1.83e+05
=====
Omnibus:      19369.210    Durbin-Watson:      1.973
Prob(Omnibus):      0.000    Jarque-Bera (JB):      1260508.231
Skew:         4.079    Prob(JB):      0.00
Kurtosis:     39.526    Cond. No.      6.37
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

This is the regression model for if the house is on the waterfront and the price. The R-squared is .071 and the F-statistic is 1650

```
In [ ]: #Supervised Regression for price and waterfront
```

```
X = cleaned_data[['waterfront']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.071
Model:                  OLS      Adj. R-squared:             0.071
Method:                 Least Squares    F-statistic:         1650.
Date:                   Sat, 13 Apr 2024    Prob (F-statistic):      0.00
Time:                   19:27:27    Log-Likelihood:        -3.0660e+05
No. Observations:      21598    AIC:                   6.132e+05
Df Residuals:          21596    BIC:                   6.132e+05
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const          5.317e+05    2416.981     219.973      0.000     5.27e+05     5.36e+05
waterfront     1.13e+06     2.78e+04     40.623      0.000     1.08e+06     1.18e+06
=====
Omnibus:                 17744.222    Durbin-Watson:           1.962
Prob(Omnibus):            0.000    Jarque-Bera (JB):        924924.672
Skew:                     3.607    Prob(JB):                 0.00
Kurtosis:                 34.237    Cond. No.                 11.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Below is the regression for if the house has a view against the price. The R-squared value is .129 and the Fstatistic is 3196.

In []: *#Supervised regression for price and view*

```
X = cleaned_data[['view']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.129
Model:                  OLS      Adj. R-squared:             0.129
Method:                 Least Squares    F-statistic:         3196.
Date:                   Sat, 13 Apr 2024    Prob (F-statistic):      0.00
Time:                   19:27:27    Log-Likelihood:        -3.0590e+05
No. Observations:      21598    AIC:                   6.118e+05
Df Residuals:          21596    BIC:                   6.118e+05
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const          4.967e+05    2455.206     202.297      0.000     4.92e+05     5.01e+05
view           4.43e+05     7832.889     56.551      0.000     4.28e+05     4.58e+05
=====
Omnibus:                 18360.975    Durbin-Watson:           1.960
Prob(Omnibus):            0.000    Jarque-Bera (JB):        1111436.469
Skew:                     3.756    Prob(JB):                 0.00
Kurtosis:                 37.331    Cond. No.                 3.40
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

This below regression goes over the relationship between price and grade. The R-squared value for this regression is .446 and the Fstatistic is 1.74.

In []: *#Supervised regression for price and grade*

```
X = cleaned_data[['grade']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())
```



```

OLS Regression Results
=====
Dep. Variable:      price      R-squared:      0.446
Model:              OLS       Adj. R-squared:    0.446
Method:             Least Squares
Date:               Sat, 13 Apr 2024
Time:              19:27:27    Prob (F-statistic): 0.00
No. Observations:   21598      Log-Likelihood:  -3.0101e+05
Df Residuals:       21596      AIC:             6.020e+05
Df Model:           1          BIC:             6.020e+05
Covariance Type:    nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const      -1.06e+06    1.23e+04    -86.368    0.000    -1.08e+06    -1.04e+06
grade       2.09e+05    1584.788    131.887    0.000     2.06e+05     2.12e+05
=====
Omnibus:            19899.052    Durbin-Watson:      1.968
Prob(Omnibus):      0.000      Jarque-Bera (JB):    2054020.764
Skew:               4.086      Prob(JB):            0.00
Kurtosis:           50.071      Cond. No.            52.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Combined Regression

Below is the summary of the regression plotted using the Ordinary Least Squares method. For this regression there was a comparison performed for all of the independent variables we use against the price.

```

In [ ]: #collection regression
X = cleaned_data[['bedrooms', 'bathrooms', 'sqft_lot', 'sqft_living', 'floors', 'waterfront', 'view', 'grade']]
Y = cleaned_data['price']

X = sm.add_constant(X)

linear_model = sm.OLS(Y, X)

results = linear_model.fit()

print(results.summary())

```

```

OLS Regression Results
=====
Dep. Variable:      price      R-squared:      0.585
Model:              OLS       Adj. R-squared:    0.585
Method:             Least Squares
Date:               Sat, 13 Apr 2024
Time:              19:27:27    Prob (F-statistic): 0.00
No. Observations:   21598      Log-Likelihood:  -2.9790e+05
Df Residuals:       21589      AIC:             5.958e+05
Df Model:           8          BIC:             5.959e+05
Covariance Type:    nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const      -5.409e+05    1.4e+04    -38.627    0.000    -5.68e+05    -5.13e+05
bedrooms    -2.469e+04    2297.994    -10.746    0.000    -2.92e+04    -2.02e+04
bathrooms   -387.2910    3318.709     -0.117    0.907    -6892.207     6117.625
sqft_lot     -0.3162         0.040     -7.825    0.000     -0.395     -0.237
sqft_living  179.0091         3.421     52.320    0.000     172.303     185.715
floors       -3.215e+04    3551.185     -9.054    0.000    -3.91e+04    -2.52e+04
waterfront   6.787e+05    1.93e+04     35.088    0.000     6.41e+05     7.17e+05
view         1.571e+05    5845.404     26.871    0.000     1.46e+05     1.69e+05
grade        1.077e+05    2248.253     47.900    0.000     1.03e+05     1.12e+05
=====
Omnibus:            16760.556    Durbin-Watson:      1.972
Prob(Omnibus):      0.000      Jarque-Bera (JB):    1324799.148
Skew:               3.139      Prob(JB):            0.00
Kurtosis:           40.851      Cond. No.            5.23e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.23e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```


In [ ]: independent_vars = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view', 'grade']
X = cleaned_data[independent_vars]

X = sm.add_constant(X)

predicted_prices = results.predict(X)

plt.figure(figsize=(10, 6))
plt.scatter(cleaned_data['price'], predicted_prices, alpha=0.5)
plt.title('Scatter Plot of Predicted vs Actual Price')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.grid(True)
plt.show()

```



Unsupervised

Below we have applied K-means clustering to compare the independent variables and the dependent variables.

This first cluster compares the sqft_living and sqft_lot with the price. The clustering is fairly spread out for this indicating the clustering wasn't very effective. It is however aparent the majority of the clustering is 2000 to 4000 sqft_living and .175 sqft_lot.

```
In [ ]: #Price with sqft_lot and sqft_living

import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'sqft_lot', 'sqft_living']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['sqft_lot'], cleaned_data['sqft_living'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('sqft_lot')
plt.ylabel('sqft_living')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



This below K-means cluster shows the relationship between price and number of bedrooms. The majority of the higher priced items and the majority of the datapoints are grouped in the middle and displays that just because there are more bedrooms, the price doesn't necessarily increase.

This cluster shows the relationship between bedrooms and price. This conflicts with the bathroom data where we don't see the number of bedrooms necessarily leading to a higher price but rather the clustered items in the 4 and 6 bedroom cluster show the highest prices and the most data as well.

```
In [ ]: #Price and bedrooms
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'bedrooms']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['bedrooms'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('bedrooms')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



This cluster shows the relationship between the price and bathrooms where the majority of the datapoints are clustered between 1 and 4.5 bathrooms with the price highest towards the upper side, indicating that there is a relationship between the number of bathrooms and the price.

```
In [ ]: #Price and bathrooms
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'bathrooms']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['bathrooms'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('bathrooms')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



This below K-means cluster shows the relationship between price and bedrooms. Based off of the layout of the cluster the majority of the houses are 2 floor houses. These houses always appear to have the highest prices. The higher and lower ends had the most diverse spread.

```
In [ ]: #price and floors
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'bedrooms']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['floors'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('floors')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



For this K-means cluster, features were clustered by price and waterfront, it doesn't seem that the clusters were

Below is a K-means cluster where the features were clustered by price and waterfront. Based off of the clustering model, it doesn't appear a relationship between waterfront and price is very strong as we see strong clustering in both waterfront properties and not.

```
In [ ]: #price and waterfront

import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

```
X = cleaned_data[['price', 'waterfront']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['waterfront'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('waterfront')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



Below is a K-means cluster that clusters the pricing and the view together. Based off this cluster it looks like there are further outliers associated with a nice view and the associated price. Indicating that the price increases with a nice view.

```
In [ ]: #Price and view
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'view']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['view'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('view')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



This is a k-means cluster that shows the relationship between price and grade. We can see lots of the clustering in yellow are all high price and also high grade.

```
In [ ]: #price and grade
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = cleaned_data[['price', 'grade']]

kmeans = KMeans(n_clusters=3)

kmeans.fit(X)

cluster_labels = kmeans.labels_

cleaned_data['cluster'] = cluster_labels

plt.figure(figsize=(10,6))


plt.scatter(cleaned_data['price'], cleaned_data['grade'], c=cleaned_data['cluster'], cmap='viridis', s=50, alpha=0.5)

plt.xlabel('price')
plt.ylabel('grade')
plt.title('Kmeans Clustering of Housing Data')

plt.colorbar(label='Cluster')

plt.show()
```

c:\Users\19noa\miniconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



All variables used for analysis: price, bedroom, bathroom, sqft_living, sqft_lot, floors, waterfront, view, grade

Train, Test, and Provide Accuracy and Evaluation Metrics

Here we train, test and provide the accuracy and evaluation metrics. The results display the R-Squared at .617 this means that the model built is responsible for 62% of variability in the house prices based on the variables included. The Mean Squared Error means that the squared difference between the predicted house prices and the actual house prices is approximately 236398.63 and for the Mean Absolute Error, on average the absolute difference between the predicted house prices and the actual house prices is approximately 154430.02.

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```

from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

X = cleaned_data[['bedrooms', 'bathrooms', 'sqft_lot', 'sqft_living', 'floors', 'waterfront', 'view', 'grade']]
Y = cleaned_data['price']

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

lm_model = LinearRegression()
lm_model.fit(X_train, y_train)

y_pred = lm_model.predict(X_test)

accuracy = r2_score(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)
mae = mean_absolute_error(y_test, y_pred)

print("R-squared: ", round(accuracy, 4))
print("Root Mean Squared Error: ", round(rmse, 4))
print("Mean absolute Error: ", round(mae, 4))

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title('Scatter Plot with Linear Regression Line')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()

```

R-squared: 0.5816
Root Mean Squared Error: 243348.4396
Mean absolute Error: 153107.9914

